

Handout 01: Statistical Learning (Optional)

Introduction

The general course this year will focus on the application of statistical learning to the analysis of textual corpora. I do understand, however, that some students may be looking for a bit more technical details than we have time to go through in the notes. To address this, I will include several handouts from a previous semester's offering of the class. But please note that these are entirely optional!

Formalisms and terminology

Here we give a formalization of the central elements of supervised learning. The resulting language and terminology will be useful as a reference.

Assume that there exists some unknown function f that maps elements from a set Ω into a set \mathcal{Y} ,

$$f : \Omega \rightarrow \mathcal{Y}, \quad (1.1)$$

and consider observing tuples, known as *training data*,

$$\{\omega_i, y_i = f(\omega_i)\}_i, \quad \omega_i \in \Omega, y_i \in \mathcal{Y}, \quad i = 1, \dots, n. \quad (1.2)$$

We will avoid writing out a formal definition of the (possibly) random nature of f and the (possibly) random process that generates each ω_i . Doing so here would overly divert from the main discussion; we will define the random nature of the data generation process whenever necessary.

A supervised learning algorithm constructs an estimate \hat{f} of the function f using the training data. The goal is to minimize errors in estimating f on new data for some loss function \mathcal{L} . When predicting a continuous response variable, such as an expected price, common loss functions include the squared or absolute difference between the observed and predicted values,

$$\mathcal{L}(\hat{f}(\omega_{new}), f(\omega_{new})) = \left| \hat{f}(\omega_{new}) - f(\omega_{new}) \right| \quad (1.3)$$

When the set of responses \mathcal{Y} is finite, as it would be when building a spam prediction algorithm, a common choice of \mathcal{L} is to measure the proportion of incorrectly labeled new observations,

$$\mathcal{L}(\hat{f}(\omega_{new}), f(\omega_{new})) = \begin{cases} 0, & \hat{f}(\omega_{new}) = f(\omega_{new}) \\ 1, & \text{otherwise} \end{cases} \quad (1.4)$$

Depending on the application, more complex metrics can be used. When evaluating a medical diagnostic algorithm, for instance, it may make more sense to weight

incorrectly missing a serious condition more heavily than incorrectly diagnosing a serious condition.

In nearly all supervised learning tasks, the training data will either be given as, or coercible to, a vector of real values. In other words, we can write the set Ω in Equation 1.1 as \mathbb{R}^p for some number p . Similarly, the prediction task can usually be re-written such that \mathcal{Y} is equal to \mathbb{R} —in the case of a discrete set, this can be done by associating each category with an integer-based index. Then, by stacking the n training inputs together, we have

$$X = \begin{pmatrix} \omega_1 \\ \vdots \\ \omega_n \end{pmatrix} \in \mathbb{R}^{n \times p}, \quad y = \begin{pmatrix} y_1 \\ \vdots \\ y_n \end{pmatrix} \in \mathbb{R}^n, \quad y = f(X). \quad (1.5)$$

The matrix X is known as the *feature matrix*. This simplification allows for us to draw on techniques from numerical analysis, functional analysis, and statistics in the pursuit of predictive models. A central theme of this course will be building and evaluating supervised learning algorithms motivated by the study of properties of the matrix X and assumptions made regarding the function f .

If the success of a supervised learning algorithm is defined on data that is, by definition, unavailable, how will it be possible to determine how well a predictive model is able to estimate the function f ? One approach is to take the observed tuples of data available for building predictive models and partition them into two subsets. Only one of these partitions is used as the *training set* to produce the estimate \hat{f} . The remaining partition, the *testing set*, is used to evaluate how well the estimate can make predictions on new data. More complex schemes operate similarly by splitting the data multiple times (e.g., *cross-validation*) or include a third partition to allow for tuning hyperparameters in the model estimation algorithm.

When considering the construction of a predictive model, a key concern is the desired capacity, or complexity, of a model building algorithm. A formal definition of a training algorithm's complexity is given by its Vapnik–Chervonenkis (VC) dimension [1], though an informal understanding of complexity will suffice here. A model that is overly complex will *overfit* to the training data; that is, \hat{f} will fit the training data very closely but not be able to generalize well to the testing data. Conversely, a model with low complexity may be too constrained to provide a good approximation to f . In a probabilistic framework, this can be seen as equivalent to a trade-off between *bias* and *variance*. A model building algorithm with high complexity will have a large variance (it may overfit, and is therefore highly sensitive to the training data); if it has a complexity that is too low to approximate f , it will provide systematically biased results for some inputs. The study and control of model complexity, in both its numerical and probabilistic forms, is a guiding theme throughout this text.

References

- [1] VAPNIK, V., AND CHERVONENKIS, A. On the uniform convergence of relative frequencies of events to their probabilities. *Theory of Probability and Its Applications* 16, 2 (1971), 264.